

Submission Guide

How to package and send in your GenAR AI Engineering Challenge.

1. What to include

```
your-name-genar-challenge/  
├── .git/                ← optional, see note below  
├── src/ (or app/, wherever your code lives)  
├── README.md  
├── architecture.png / .pdf / .md    ← your diagram  
├── report_output.md / .html / .docx / .pdf ← the generated PADER-style report  
├── prompts/ (or inline in code – just make them easy to find)  
└── version1/ ← implementation, OR a one-page design doc if you didn't build it
```

You don't need to match this layout exactly — this is a checklist of what should exist somewhere in the zip, not a required folder structure.

Do not include: `venv/`, `node_modules/`, `__pycache__/`, `.env`, model caches, or the dataset CSV itself (we already have it — re-including it just eats your size budget).

2. `.git` history (optional)

If you happened to commit along the way, feel free to leave `.git/` in the zip rather than stripping it out — it's not required, and a single "final" commit or no repo at all isn't held against you. It's just that if you *did* work incrementally, the history is an easy way for it to speak for itself at the end, on top of whatever you say in the README.

3. Size limit

Target $\leq 30\text{MB}$ zipped. If you're over that, it's almost always one of: a virtual env or `node_modules` that snuck in, an embedded copy of the dataset, or committed model weights/caches. None of those should be in the zip.

4. README — what it needs to answer

Not a template to fill in mechanically — but by the end, a reader should be able to answer all of these from your README alone:

- **How do I run it?** (setup + one command to regenerate the report)

- **What's the architecture?** (matches your diagram — where does data go, where does AI get involved, where doesn't it)
- **Where is AI used vs. deterministic code, and why that split?**
- **What are the actual prompts/context templates your system assembles?** (show them, don't just describe them)
- **How does the system stay grounded?** (i.e., how do you know a generated sentence is backed by the data)
- **How would you evaluate this at scale** (1,000 generated reports, not one)?
- **What are the known limitations?** (this one matters — "I didn't handle X, here's how I would" is a good answer, silence on obvious gaps is not)

5. Architecture diagram

Any format is fine — hand-drawn and photographed, Excalidraw export, a Mermaid block in your README, whatever. It needs to show components and data flow, not be polished.

6. Heads-up: expect a live follow-up

After we look at your submission, expect a short live session where we ask you to **walk through a specific design decision** and **extend or fix something in your own system on the spot**. Build accordingly — you should understand every part of what you submit well enough to explain and change it live, regardless of what tools you used to build it.

7. Where to send it

Follow the communication.

Name the zip `firstname_lastname_genar_challenge.zip` .

That's it. If you're unsure whether something belongs, err on the side of including it in the README as a decision you made rather than leaving us to guess.